

Least Squares / Linear Regression

Sometimes we're given some data points of which we need to find a trend line or otherwise known as a regression line.

Let's say we're given a couple points, (2, 3) and (7, 4). The points really represent two equations (where we automatically add a coefficient of 1):

$$f(2) = \beta_0 + \beta_1 2 = 3$$

$$f(7) = \beta_0 + \beta_1 7 = 4$$

$$\begin{pmatrix} 1 & 2 \\ 1 & 7 \end{pmatrix} \begin{pmatrix} \beta_0 \\ \beta_1 \end{pmatrix} = \begin{pmatrix} 3 \\ 4 \end{pmatrix}$$

A \vec{x} \vec{b}

If we take the matrix A, augment it with \vec{b} and row-reduce the entire thing, you'll find we get:

$$\left(\begin{array}{cc|c} 1 & 0 & \frac{13}{5} \\ 0 & 1 & \frac{1}{5} \end{array} \right)$$

Now we know our beta parameters:

$$\beta_0 = \frac{13}{5}$$

$$\beta_1 = \frac{1}{5}$$

Ok, that's all cool and all, but sometimes we'll be given some data where after row-reducing, doesn't have a solution. The below is an example of a row-reduced matrix.

$$\left(\begin{array}{cc|c} 1 & 0 & \frac{13}{5} \\ 0 & 1 & \frac{1}{5} \\ 0 & 0 & \frac{18}{5} \end{array} \right)$$

You can see that the bottom row is inconsistent, $0 \neq \frac{18}{5}$. No line fits these points so we need to find a line that best fits these points, i.e. we need an approximate solution. This leads us to the Least Squares Solution, $A\hat{x} = \vec{b}^{\perp}$.

The Long Way

Let's first approach this the long way.

Say you have the following points: (2,3),(7,4),(9,8).

Since \vec{b} is not in the range(A), we need to project it onto the range(A). Therefore, $A\hat{x} = \vec{b}^{\perp}$ has a solution even though $A\vec{x} = \vec{b}$ does not (which commonly happens when $m > n$). Why \vec{b}^{\perp} ? Because it's the closest vector in range(A) to \vec{b} (The Best Approximation Theorem), which means whatever coefficients we get for \vec{x} is really \hat{x} because they're approximated.

The algo:

1. Find A. It's columns are $x[1]$, $x[2]$ below. $b = \{3, 4, 8\}$, the "observation" vector.
2. Gram-Schmidt-ify to get an orthogonal basis - $v[1], v[2]$.
3. Find \vec{b}^{\perp} , the orthogonal projection of \vec{b} onto span of $v[1], v[2]$.
4. Solve $A\hat{x} = \vec{b}^{\perp}$.

$$\begin{pmatrix} 1 & 2 \\ 1 & 7 \\ 1 & 9 \end{pmatrix} * \begin{pmatrix} \beta_0 \\ \beta_1 \end{pmatrix} = \begin{pmatrix} 3 \\ 4 \\ 8 \end{pmatrix}$$

```
Clear[b, x, proj, v, u]
b = {3, 4, 8};
x[1] = {1, 1, 1};
x[2] = {2, 7, 9};
p = 2; (*number of parameters*)
A = {x[1], x[2]}^T;
Print["A=", A // MatrixForm]
Print["b=", b // MatrixForm]
```

$$A = \begin{pmatrix} 1 & 2 \\ 1 & 7 \\ 1 & 9 \end{pmatrix}$$

$$b = \begin{pmatrix} 3 \\ 4 \\ 8 \end{pmatrix}$$

Gram-Schmidt-ify to get an orthogonal basis.

```

For[i = 1, i ≤ p, i++,
  proj[i] =  $\sum_{j=1}^{i-1} \left( \frac{\mathbf{x}[i] \cdot \mathbf{v}[j]}{\mathbf{v}[j] \cdot \mathbf{v}[j]} \right) \mathbf{v}[j];$ 
  v[i] = x[i] - proj[i];
  u[i] =  $\frac{\mathbf{v}[i]}{\sqrt{\mathbf{v}[i] \cdot \mathbf{v}[i]}}$ ;
]
proj[1] = x[1] * 0;
T = Table[{x[i] // MatrixForm, proj[i] // Expand // MatrixForm,
  v[i] // Expand // MatrixForm, u[i] // Expand // MatrixForm}, {i, 1, p}];
T = Prepend[T, {"original vector",
  "orthogonal projection onto subspace spanned by prior vectors",
  "New Gram-Schmidt-ified vector", "Normalized"}];
T = Prepend[T, {"xi", " $\mathbf{x}_i'' = \sum_{j=1}^{i-1} \left( \frac{\mathbf{x}_i \cdot \mathbf{v}_j}{\mathbf{v}_j \cdot \mathbf{v}_j} \right) \mathbf{v}_j$ ", " $\mathbf{v}_i = \mathbf{x}_i - \mathbf{v}_i''$ ", " $\mathbf{u}_i = \frac{\mathbf{v}_i}{\|\mathbf{v}_i\|}$ "}];
Print[T // TableForm]

```

\mathbf{x}_i	$\mathbf{x}_i'' = \sum_{j=1}^{i-1} \left(\frac{\mathbf{x}_i \cdot \mathbf{v}_j}{\mathbf{v}_j \cdot \mathbf{v}_j} \right) \mathbf{v}_j$	$\mathbf{v}_i = \mathbf{x}_i - \mathbf{v}_i''$
original vector	orthogonal projection onto subspace spanned by prior vectors	New Gram
$\begin{pmatrix} 1 \\ 1 \\ 1 \end{pmatrix}$	$\begin{pmatrix} 0 \\ 0 \\ 0 \end{pmatrix}$	$\begin{pmatrix} 1 \\ 1 \\ 1 \end{pmatrix}$
$\begin{pmatrix} 2 \\ 7 \\ 9 \end{pmatrix}$	$\begin{pmatrix} 6 \\ 6 \\ 6 \end{pmatrix}$	$\begin{pmatrix} -4 \\ 1 \\ 3 \end{pmatrix}$

Find $\vec{\hat{b}}$.

```

For[i = 1, i ≤ p, i++,
  c[i] =  $\frac{\mathbf{b} \cdot \mathbf{v}[i]}{\mathbf{v}[i] \cdot \mathbf{v}[i]}$ ;
  Print["c[" , i, "] = ", c[i]]
]
bp =  $\sum_{i=1}^p \mathbf{c}[i] * \mathbf{v}[i];$ 
Print["b" = ", bp // MatrixForm]

```

c[1]=5

$c[2] = \frac{8}{13}$

$\mathbf{b}'' = \begin{pmatrix} \frac{33}{13} \\ \frac{73}{13} \\ \frac{89}{13} \end{pmatrix}$

Solve $A\hat{\mathbf{x}} = \vec{\hat{b}}$.

```

B = Append[AT, bp]T;
Print["B=[A|b]=", B // MatrixForm];
R = RowReduce[B];
Print["R=rref[B]=", R // MatrixForm];

```

$$B = [A|b] = \begin{pmatrix} 1 & 2 & \frac{33}{13} \\ 1 & 7 & \frac{73}{13} \\ 1 & 9 & \frac{89}{13} \end{pmatrix}$$

$$R = \text{rref}[B] = \begin{pmatrix} 1 & 0 & \frac{17}{13} \\ 0 & 1 & \frac{8}{13} \\ 0 & 0 & 0 \end{pmatrix}$$

The Short Way

Now we can use a faster way to solve $A\hat{x} = \vec{b}$, using Thm 6.58.13 (the Normal Equation).

Remember that the solutions to $A\hat{x} = \vec{b}$ are also the solutions to $A^T A\hat{x} = A^T \vec{b}$.

```

AT.A // MatrixForm
Inverse[AT.A] // MatrixForm
Inverse[AT.A].AT // MatrixForm
β = Inverse[AT.A].AT.b;
Print["Least squares solution = ", β // MatrixForm]
e = b - A.β;
Print["Residual vector = ", e // MatrixForm]
Print["Residual = ", Norm[e] // N]

```

$$\begin{pmatrix} 3 & 18 \\ 18 & 134 \end{pmatrix}$$

$$\begin{pmatrix} \frac{67}{39} & -\frac{3}{13} \\ -\frac{3}{13} & \frac{1}{26} \end{pmatrix}$$

$$\begin{pmatrix} \frac{49}{39} & \frac{4}{39} & -\frac{14}{39} \\ -\frac{2}{13} & \frac{1}{26} & \frac{3}{26} \end{pmatrix}$$

$$\text{Least squares solution} = \begin{pmatrix} \frac{17}{13} \\ \frac{8}{13} \end{pmatrix}$$

$$\text{Residual vector} = \begin{pmatrix} \frac{6}{13} \\ -\frac{21}{13} \\ \frac{15}{13} \end{pmatrix}$$

$$\text{Residual} = 2.0381$$

Now a new example with randomly generated data.

```

(*generate random data*)
n = 30;
M = 10;
a = RandomReal[{0, M}, n];
g = 2 - 5 t + 40 t^2;
b = g /. {t -> a};
e = RandomVariate[NormalDistribution[0, 100], n];
b = b + e;
ListPlot[Transpose[{a, b}], PlotStyle -> PointSize[Large]]

(*Compute Correlation coefficient*)
x = a - Mean[a];
y = b - Mean[b];
r =  $\frac{x \cdot y}{\text{Norm}[x] \cdot \text{Norm}[y]}$ ;
Print["Correlation coefficient is r = ", N[r]]

(*Establish model*)
Clear[t];
f = {1, t};

(*Construct A*)
A = Table[f /. t -> a[[i]], {i, 1, Length[a]}];
Print["Model function is f = ", f]
Print["A = ", A // MatrixForm];


(*Compute and plot least squares fit*)
 $\beta$  = Inverse[AT.A].AT.b;
Print["Least squares solution = ",  $\beta$  // MatrixForm]
e = b - A. $\beta$ ;
Print["Residual = ", Norm[e] // N]
h = f. $\beta$ ;
Print["Best fit model function = ", h]
Show[ListPlot[Transpose[{a, b}], PlotStyle -> PointSize[Large]],
  Plot[h, {t, Min[a] - 1, Max[a] + 1}]]

(*Establish model*)
Clear[t];
f = {1, t, t^2};

(*Construct A*)
A = Table[f /. t -> a[[i]], {i, 1, Length[a]}];
Print["Model function is f = ", f]
Print["A = ", A // MatrixForm];

(*Compute and plot least squares fit*)
 $\beta$  = Inverse[AT.A].AT.b;
Print["Least squares solution = ",  $\beta$  // MatrixForm]
e = b - A. $\beta$ ;
Print["Residual = ", Norm[e] // N]
h = f. $\beta$ ;
Print["Best fit model function = ", h]
Show[ListPlot[Transpose[{a, b}], PlotStyle -> PointSize[Large]],
  Plot[h, {t, Min[a] - 1, Max[a] + 1}]]

```



Correlation coefficient is $r = 0.970527$

Model function is $f = \{1, t\}$


$A =$

1	1.1996
1	7.7576
1	2.13487
1	8.05847
1	7.63219
1	0.0513538
1	5.36849
1	8.77227
1	2.01908
1	0.64289
1	6.35248
1	7.569
1	2.43752
1	4.10831
1	6.615
1	1.59084
1	2.17268
1	1.11084
1	6.68779
1	5.73913
1	8.07761
1	3.27042
1	3.45693
1	7.64491
1	4.57847
1	8.599
1	1.72204
1	9.86106
1	6.68554
1	5.43943

Least squares solution = $\begin{pmatrix} -599.841 \\ 379.512 \end{pmatrix}$

Residual = 1465.83

Best fit model function = $-599.841 + 379.512 t$



Model function is $f = \{1, t, t^2\}$


$A =$

1	1.1996	1.43904
1	7.7576	60.1803
1	2.13487	4.55769
1	8.05847	64.9389
1	7.63219	58.2503
1	0.0513538	0.00263722
1	5.36849	28.8207
1	8.77227	76.9528
1	2.01908	4.07669
1	0.64289	0.413308
1	6.35248	40.354
1	7.569	57.2897
1	2.43752	5.94152
1	4.10831	16.8782
1	6.615	43.7582
1	1.59084	2.53076
1	2.17268	4.72054
1	1.11084	1.23396
1	6.68779	44.7266
1	5.73913	32.9376
1	8.07761	65.2478
1	3.27042	10.6956
1	3.45693	11.9504
1	7.64491	58.4446
1	4.57847	20.9624
1	8.599	73.9428
1	1.72204	2.96543
1	9.86106	97.2405
1	6.68554	44.6965
1	5.43943	29.5874

Least squares solution = $\begin{pmatrix} -15.4032 \\ 1.37295 \\ 39.5426 \end{pmatrix}$

Residual = 447.123

Best fit model function = $-15.4032 + 1.37295 t + 39.5426 t^2$



More Examples

```

(*generate random data*)
n = 50;
M = 100;
a = RandomReal[{0, M}, n];
g = -9.8 t^2 + 950 t;
b = g /. {t → a};
e = RandomVariate[NormalDistribution[0, 1000], n];
b = b + e;
ListPlot[Transpose[{a, b}], PlotStyle → PointSize[Large]]

(*Compute Correlation coefficient*)
x = a - Mean[a];
y = b - Mean[b];
r =  $\frac{x \cdot y}{\text{Norm}[x] \cdot \text{Norm}[y]}$ ;
Print["Correlation coefficient is r = ", N[r]]

(*Establish model*)
Clear[t];
f = {1, t};

(*Construct A*)
A = Table[f /. t → a[[i]], {i, 1, Length[a]}];
Print["Model function is f = ", f]
(*Print["A = ", A//MatrixForm];*)


(*Compute and plot least squares fit*)
β = Inverse[AT.A].AT.b;
Print["Least squares solution = ", β // MatrixForm]
e = b - A.β;
Print["Residual = ", Norm[e] // N]
h = f.β;
Print["Best fit model function = ", h]
Show[ListPlot[Transpose[{a, b}], PlotStyle → PointSize[Large]],
Plot[h, {t, Min[a] - 1, Max[a] + 1}]]

(*Establish model*)
Clear[t];
f = {1, t, t2};

(*Construct A*)
A = Table[f /. t → a[[i]], {i, 1, Length[a]}];
Print["Model function is f = ", f]
(*Print["A = ", A//MatrixForm];*)

(*Compute and plot least squares fit*)
β = Inverse[AT.A].AT.b;
Print["Least squares solution = ", β // MatrixForm]
e = b - A.β;
Print["Residual = ", Norm[e] // N]
h = f.β;
Print["Best fit model function = ", h]
Show[ListPlot[Transpose[{a, b}], PlotStyle → PointSize[Large]],
Plot[h, {t, Min[a] - 1, Max[a] + 1}]]

```




Correlation coefficient is $r = -0.254976$

Model function is $f = \{1, t\}$

Least squares solution = $\begin{pmatrix} 17\,094.3 \\ -57.5037 \end{pmatrix}$

Residual = 46 728.3

Best fit model function = $17\,094.3 - 57.5037 t$




Model function is $f = \{1, t, t^2\}$

Least squares solution = $\begin{pmatrix} -78.6681 \\ 974.294 \\ -10.0404 \end{pmatrix}$

Residual = 6682.58

Best fit model function = $-78.6681 + 974.294 t - 10.0404 t^2$



Yet another example with randomly generated data.

```

(*generate random data*)
n = 85;
M = 10;
a = RandomReal[{0, M}, n];
g = Cos[t] + 2 Sin[t];
b = g /. {t -> a};
e = RandomVariate[NormalDistribution[0, .3], n];
b = b + e;
ListPlot[Transpose[{a, b}], PlotStyle -> PointSize[Large]]

(*Compute Correlation coefficient*)
x = a - Mean[a];
y = b - Mean[b];
r =  $\frac{x \cdot y}{\text{Norm}[x] \cdot \text{Norm}[y]}$ ;
Print["Correlation coefficient is r = ", N[r]]

(*Establish model*)
Clear[t];
f = {1, t};

(*Construct A*)
A = Table[f /. t -> a[[i]], {i, 1, Length[a]}];
Print["Model function is f = ", f]
(*Print["A = ", A//MatrixForm];*)


(*Compute and plot least squares fit*)
 $\beta$  = Inverse[AT.A].AT.b;
Print["Least squares solution = ",  $\beta$  // MatrixForm]
e = b - A. $\beta$ ;
Print["Residual = ", Norm[e] // N]
h = f. $\beta$ ;
Print["Best fit model function = ", h]
Show[ListPlot[Transpose[{a, b}], PlotStyle -> PointSize[Large]],
Plot[h, {t, Min[a] - 1, Max[a] + 1}]]

(*Establish model*)
Clear[t];
f = {1, Sin[t], Cos[t]};

(*Construct A*)
A = Table[f /. t -> a[[i]], {i, 1, Length[a]}];
Print["Model function is f = ", f]
(*Print["A = ", A//MatrixForm];*)

(*Compute and plot least squares fit*)
 $\beta$  = Inverse[AT.A].AT.b;
Print["Least squares solution = ",  $\beta$  // MatrixForm]
e = b - A. $\beta$ ;
Print["Residual = ", Norm[e] // N]
h = f. $\beta$ ;
Print["Best fit model function = ", h]
Show[ListPlot[Transpose[{a, b}], PlotStyle -> PointSize[Large]],
Plot[h, {t, Min[a] - 1, Max[a] + 1}]]

```




Correlation coefficient is $r = -0.216133$

Model function is $f = \{1, t\}$

Least squares solution = $\begin{pmatrix} 0.643567 \\ -0.119277 \end{pmatrix}$

Residual = 14.4434

Best fit model function = $0.643567 - 0.119277 t$




Model function is $f = \{1, \text{Sin}[t], \text{Cos}[t]\}$

Least squares solution = $\begin{pmatrix} -0.00117609 \\ 2.01716 \\ 1.05049 \end{pmatrix}$

Residual = 2.74365

Best fit model function = $-0.00117609 + 1.05049 \text{Cos}[t] + 2.01716 \text{Sin}[t]$



Example with non-random data - problem #10 in Section 6.6 (p.376).

```
a = {10, 11, 12, 14, 15};
```

```
b = {21.34, 20.68, 20.05, 18.87, 18.30};
```

```
(*Establish model*)
```

```
Clear[t];
```

```
f = {e-0.02 t, e-0.07 t};
```

```
(*Construct A*)
```

```
A = Table[f /. t → a[[i]], {i, 1, Length[a]}];
```

```
Print["Model function is f = ", f]
```

```
(*Print["A = ", A//MatrixForm];*)
```

```
(*Compute and plot least squares fit*)
```

```
 $\beta$  = Inverse[AT.A].AT.b;
```

```
Print["Least squares solution = ",  $\beta$  // MatrixForm]
```

```
e = b - A. $\beta$ ;
```

```
Print["Residual = ", Norm[e] // N]
```

```
h = f. $\beta$ ;
```

```
Print["Best fit model function = ", h]
```

```
Show[ListPlot[Transpose[{a, b}], PlotStyle → PointSize[Large]],
```


```
Plot[h, {t, Min[a] - 1, Max[a] + 1}]]
```

Model function is $f = \{e^{-0.02t}, e^{-0.07t}\}$

Least squares solution = $\begin{pmatrix} 19.9411 \\ 10.1015 \end{pmatrix}$

Residual = 0.0114745

Best fit model function = $10.1015 e^{-0.07t} + 19.9411 e^{-0.02t}$



Example with non-random data - problem #10 in Section 6.6 (p.376). We get a slightly different answer than the solution manual because their approach has round off error.

```
a = {44, 61, 81, 113, 131};
b = {91, 98, 103, 110, 112};
```

```
(*Establish model*)
Clear[t];
f = {1, Log[t]};
```

```
(*Construct A*)
A = Table[f /. t -> a[[i]], {i, 1, Length[a]}];
Print["Model function is f = ", f]
(*Print["A = ", A//MatrixForm];*)
```

```
(*Compute and plot least squares fit*)
β = Inverse[AT.A].AT.b // N;
Print["Least squares solution = ", β // MatrixForm]
e = b - A.β;
Print["Residual = ", Norm[e] // N]
h = f.β;
Print["Best fit model function = ", h]
Show[ListPlot[Transpose[{a, b}], PlotStyle -> PointSize[Large]],
Plot[h, {t, Min[a] - 1, Max[a] + 1}]]
```


Model function is $f = \{1, \text{Log}[t]\}$

Least squares solution = $\begin{pmatrix} 17.9243 \\ 19.385 \end{pmatrix}$

Residual = 0.784091

Best fit model function = $17.9243 + 19.385 \text{Log}[t]$

